

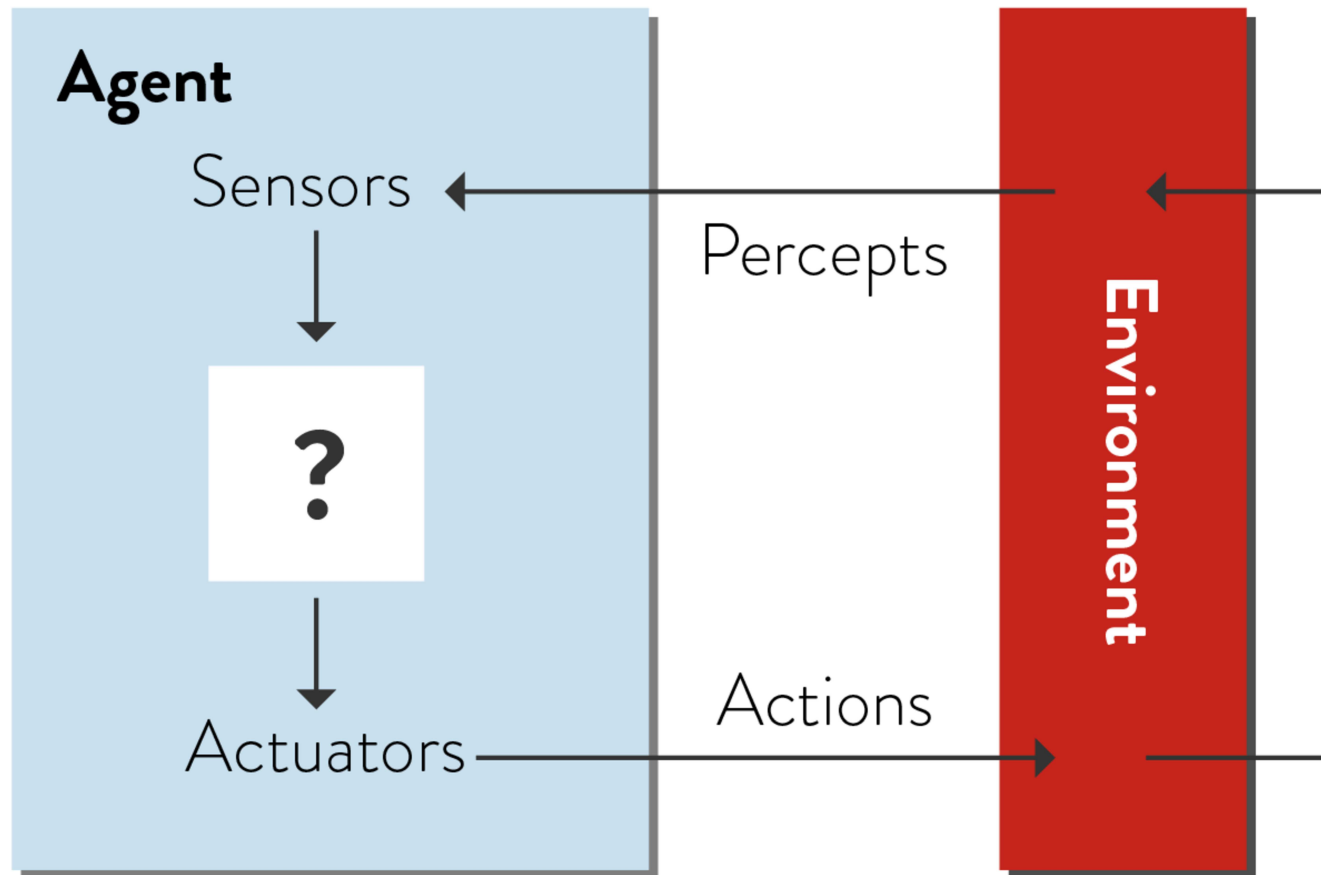
# Agents

**COMP3411/9814:  
Artificial Intelligence**

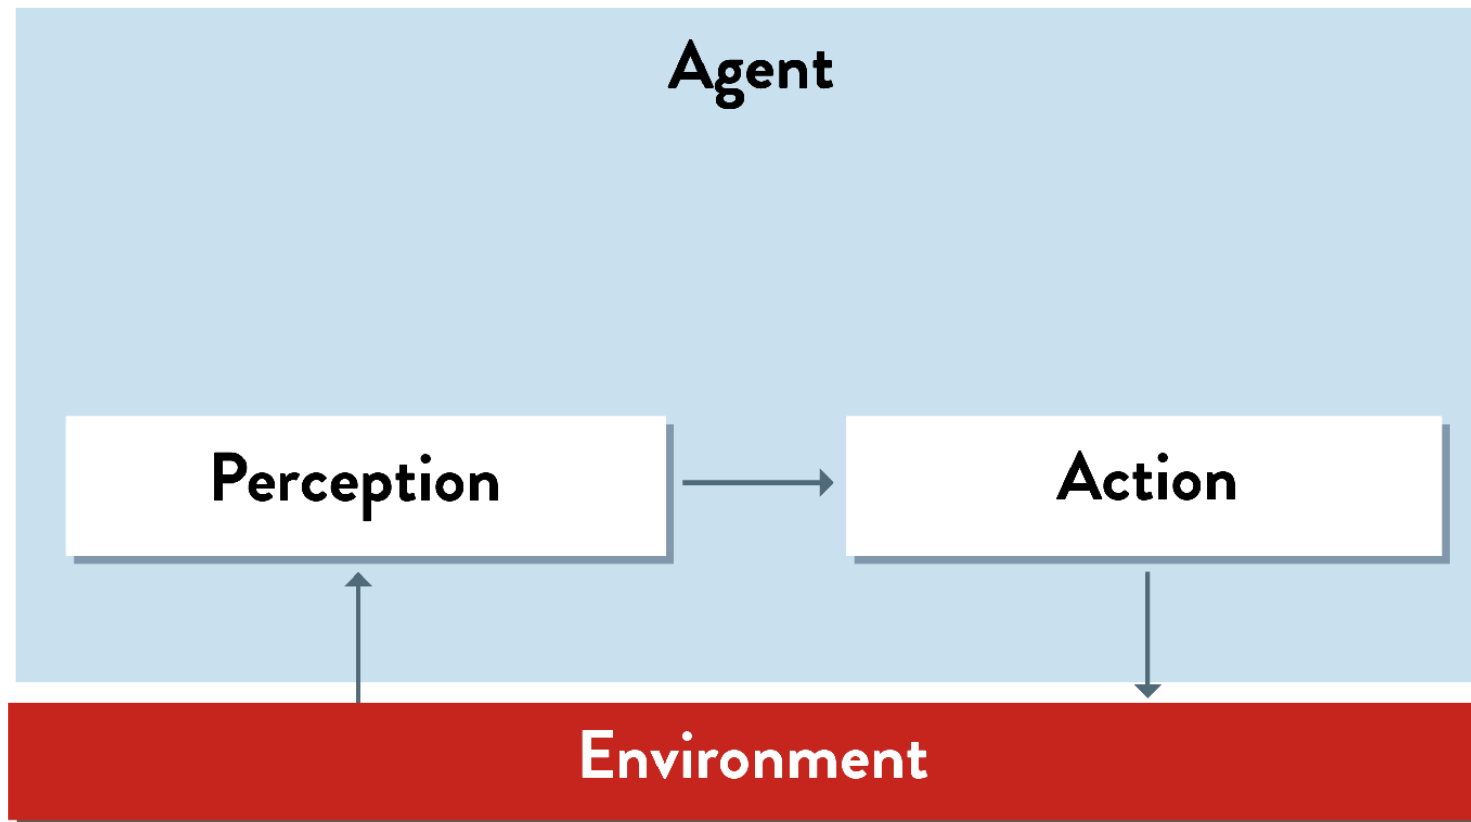
# Types of Agents

- Reactive Agent
- Model-Based Agent
- Planning Agent
- Utility-based agent
- Game Playing Agent
- Learning Agent

# Agent Model



# Reactive Agent



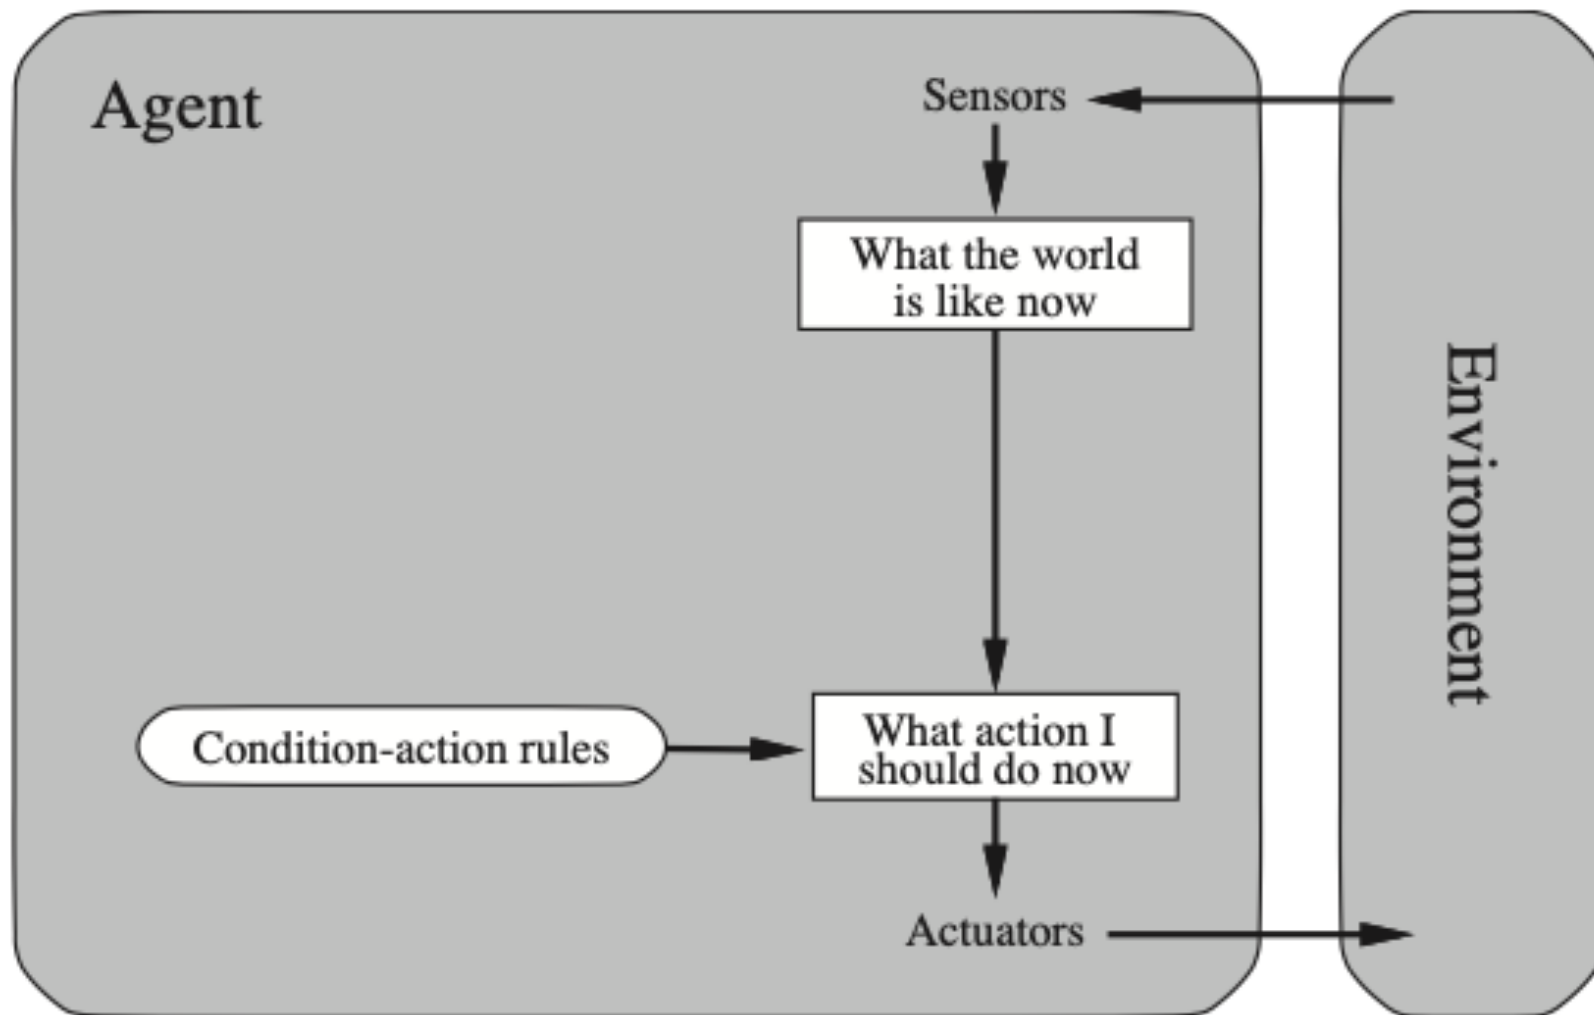
# Reactive Agent

- Choose the next action based only on what agent currently perceives
  - Uses a “policy” or set of rules that are simple to apply
- Sometimes called “simple reflex agents”
  - but they can do surprisingly sophisticated things

# Reactive Agent



# Reflex (Reactive) Agent




Reflex (reactive) agent — applies condition-action rules to each percept

# Reactive Robots






# Vacuum-cleaner world



```
function REFLEX-VACUUM-AGENT([location,status]) returns an action
  if status = Dirty then return Suck
  else if location = A then return Right
  else if location = B then return Left
```

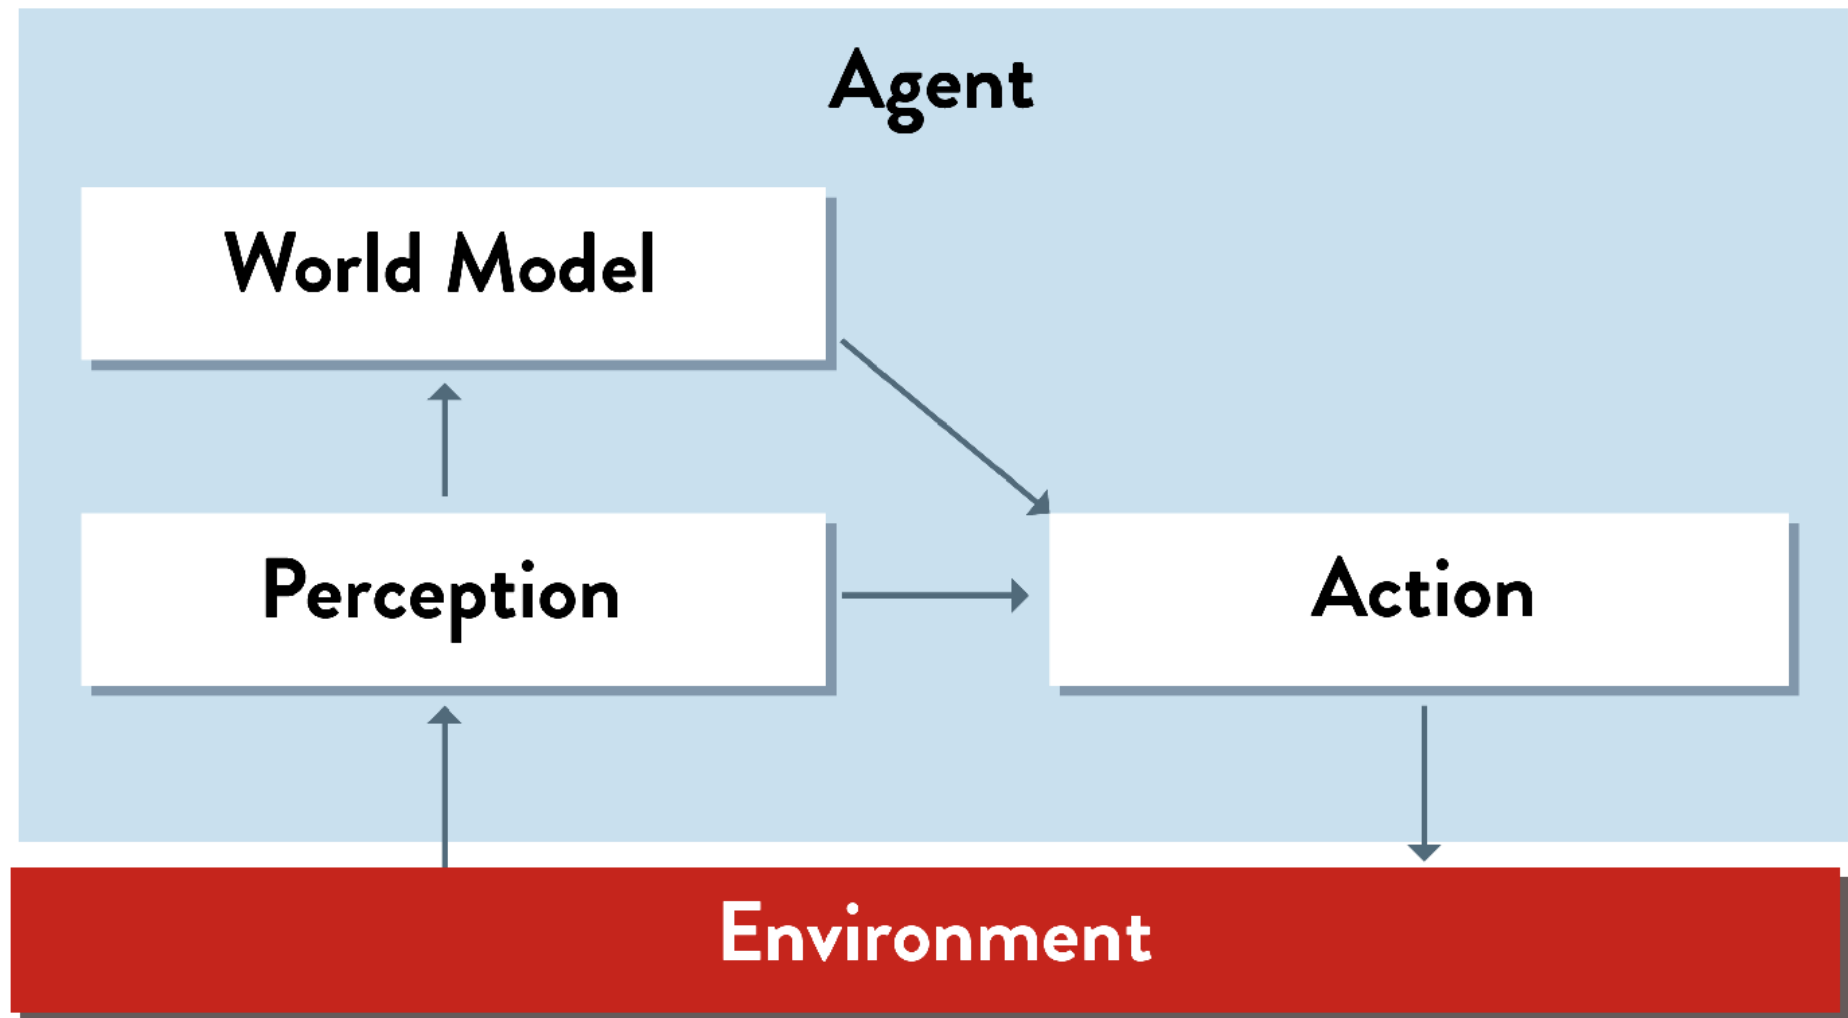
# Limitations of Reactive Agents



# Limitations of Reactive Agents

- Reactive Agents have no memory or “state”
  - unable to base decision on previous observations
  - may repeat the same sequence of actions over and over
  - Escape from infinite loops is possible if the agent can randomise its actions.

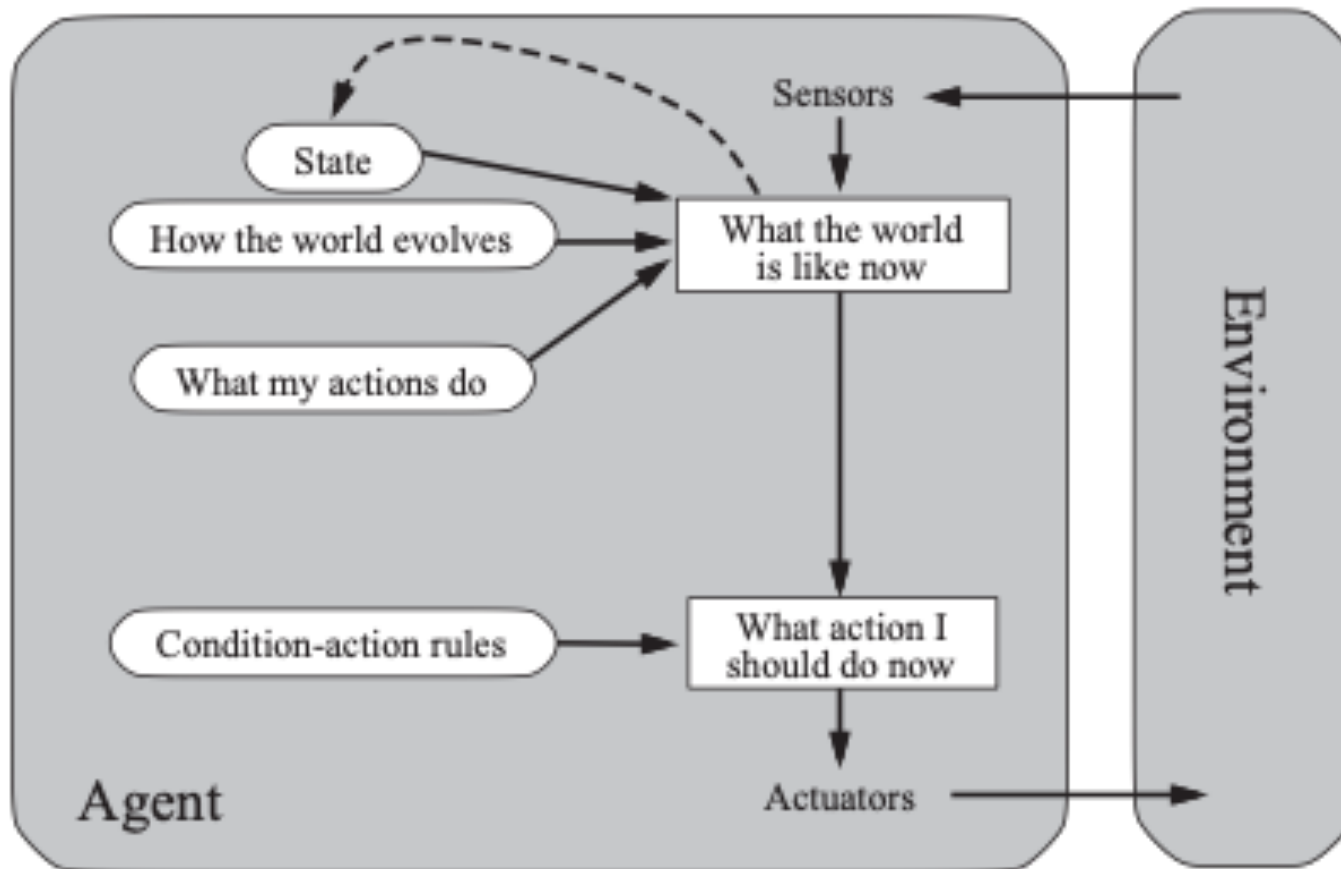
# Model-Based Agent



# Model-based Agents


- Handle *partial observability* by *keeping track of the part of the world it can't see now*.
- Maintain internal state that depends on the percept history and remembers at least some of the unobserved aspects of the current state.
- Knowledge about “how the world works” is called a **model** of the world.
- An agent that uses such a model is called a **model-based agent**.

# Model-based Reflex Agent



A model-based reflex agent. It keeps track of the current state of the world, using an internal model. It then chooses an action in the same way as the reflex agent.

# Model-based Reflex Agent



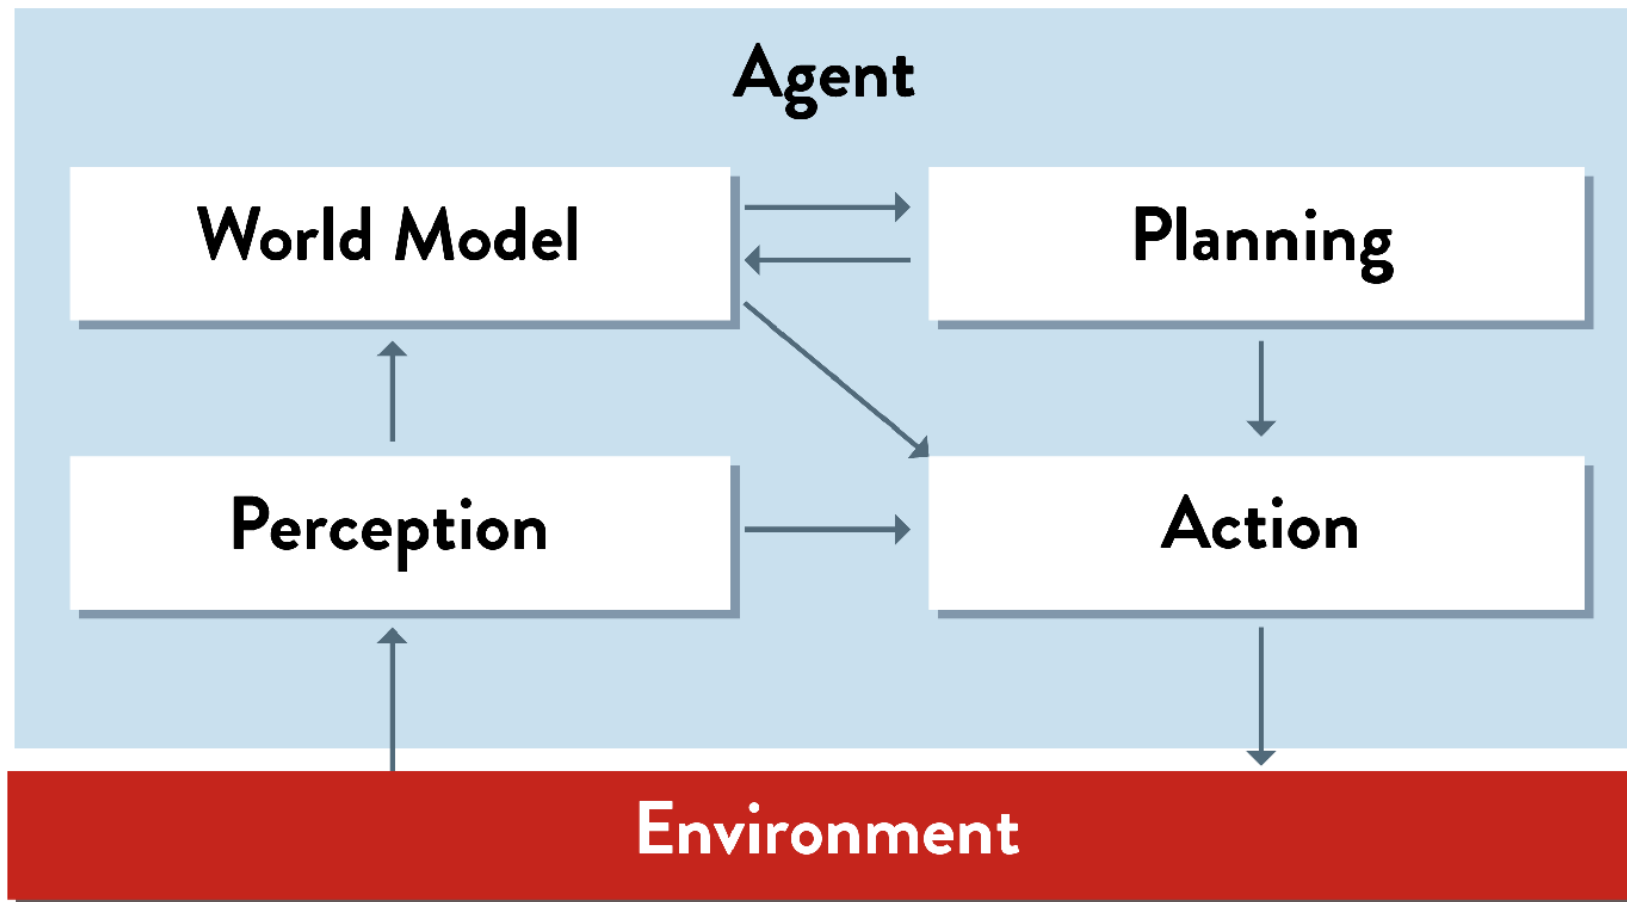
# Limitations of Model-Based Agents

- An agent with a world model but no planning can look into the past, but not into the future; it will perform poorly when the task requires any of the following:
  - searching several moves ahead
    - Chess, Rubik's cube
  - complex tasks requiring many individual step
    - cooking a meal, assembling a watch
  - logical reasoning to achieve goals
    - travel to New York

Sometimes we may need to plan several steps into the future



# Planning Agent



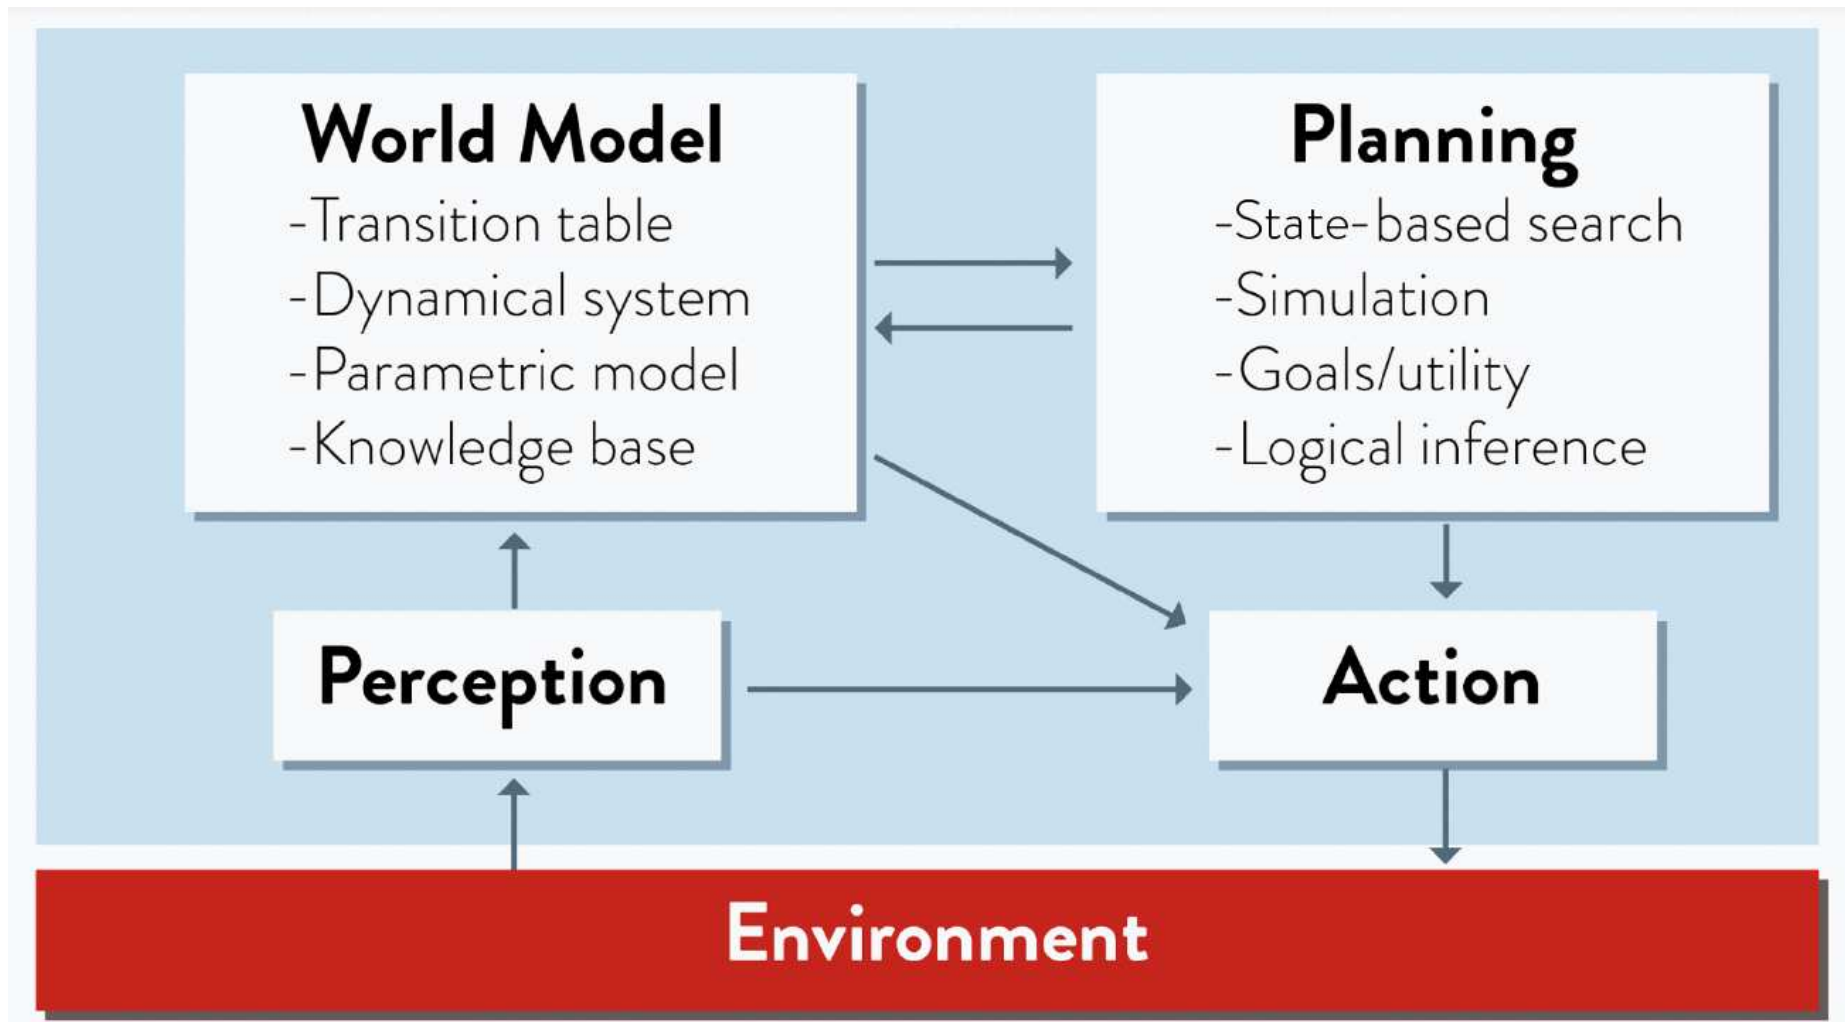
Goal-Based Agent

# Planning Agent

- Decision making of this kind is fundamentally different from the condition-action rules
- It involves consideration of the future
  - “What will happen if I do such-and-such?” and
  - “Will that make me happy?”

In the reflex agent designs, this information is not explicitly represented, because the built-in rules map directly from states to actions

# Models and Planning



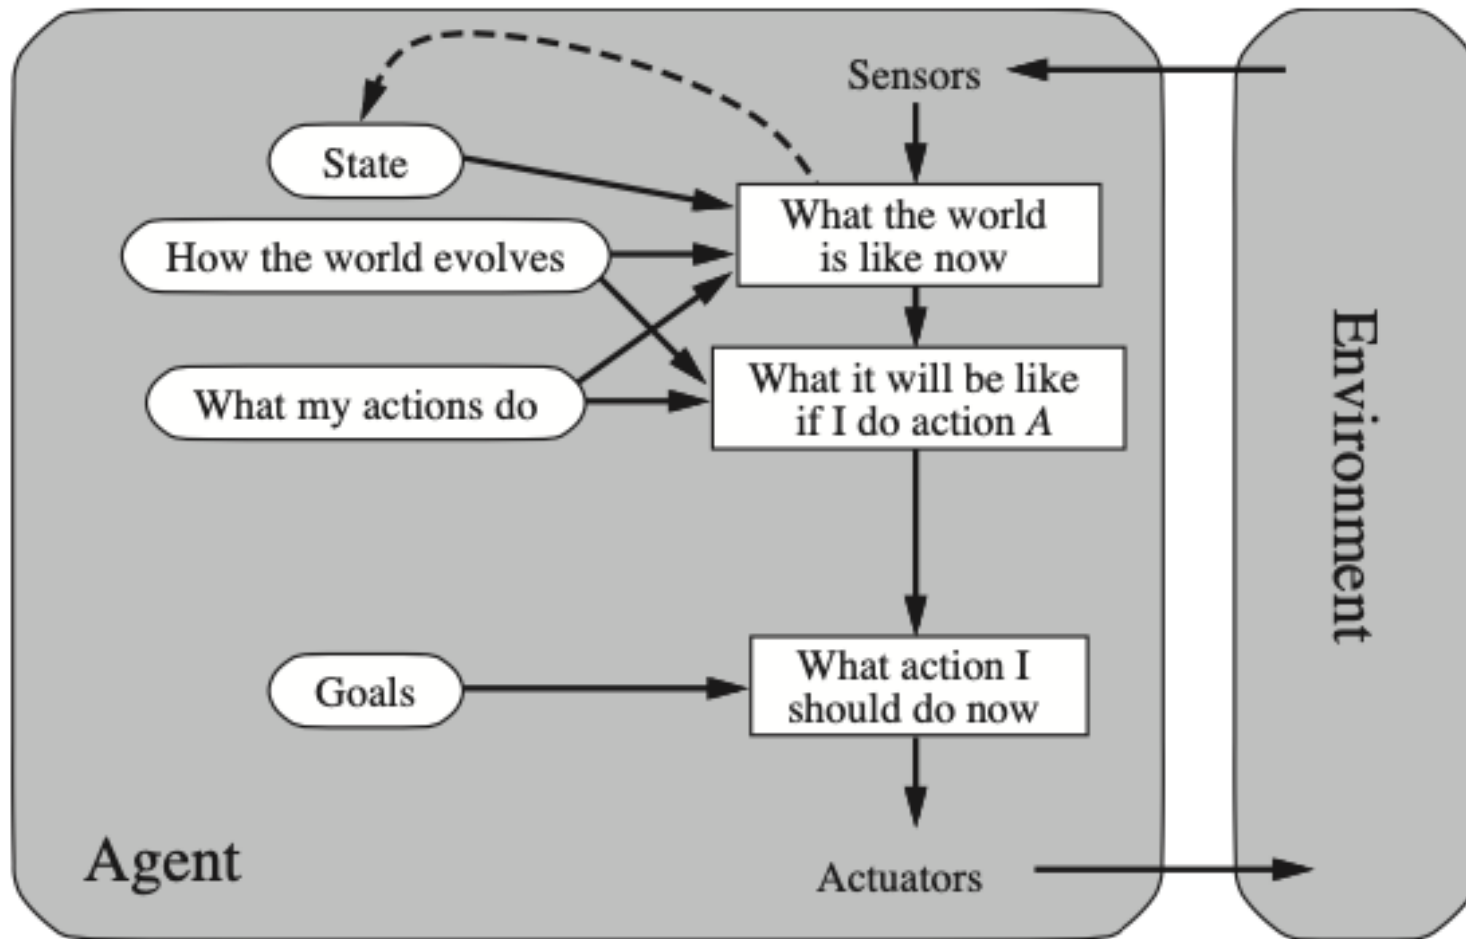
# Reasoning about Future States

- What is the best action in this situation?
- Faking it
  - Sometimes an agent may appear to be planning ahead but is actually just applying reactive rules.
  - These rules can be hand-coded, or learned from experience.
  - Agent may not be flexible enough to adapt to new situations.

# Planning Agent – Goal-based

- The planning agent or goal-based agent is more flexible because the knowledge that supports its decisions is represented explicitly and **can be modified**.
- The agent's behaviour can easily be changed.
- But ...
  - **it's slower to react because it has to “think” about what it's doing.**

# Goal-based (teleological) agent



Goal-based (teleological) agent — state description often not sufficient for agent to decide what to do so it needs to consider its goals (may involve searching and planning)

# @Home Robot

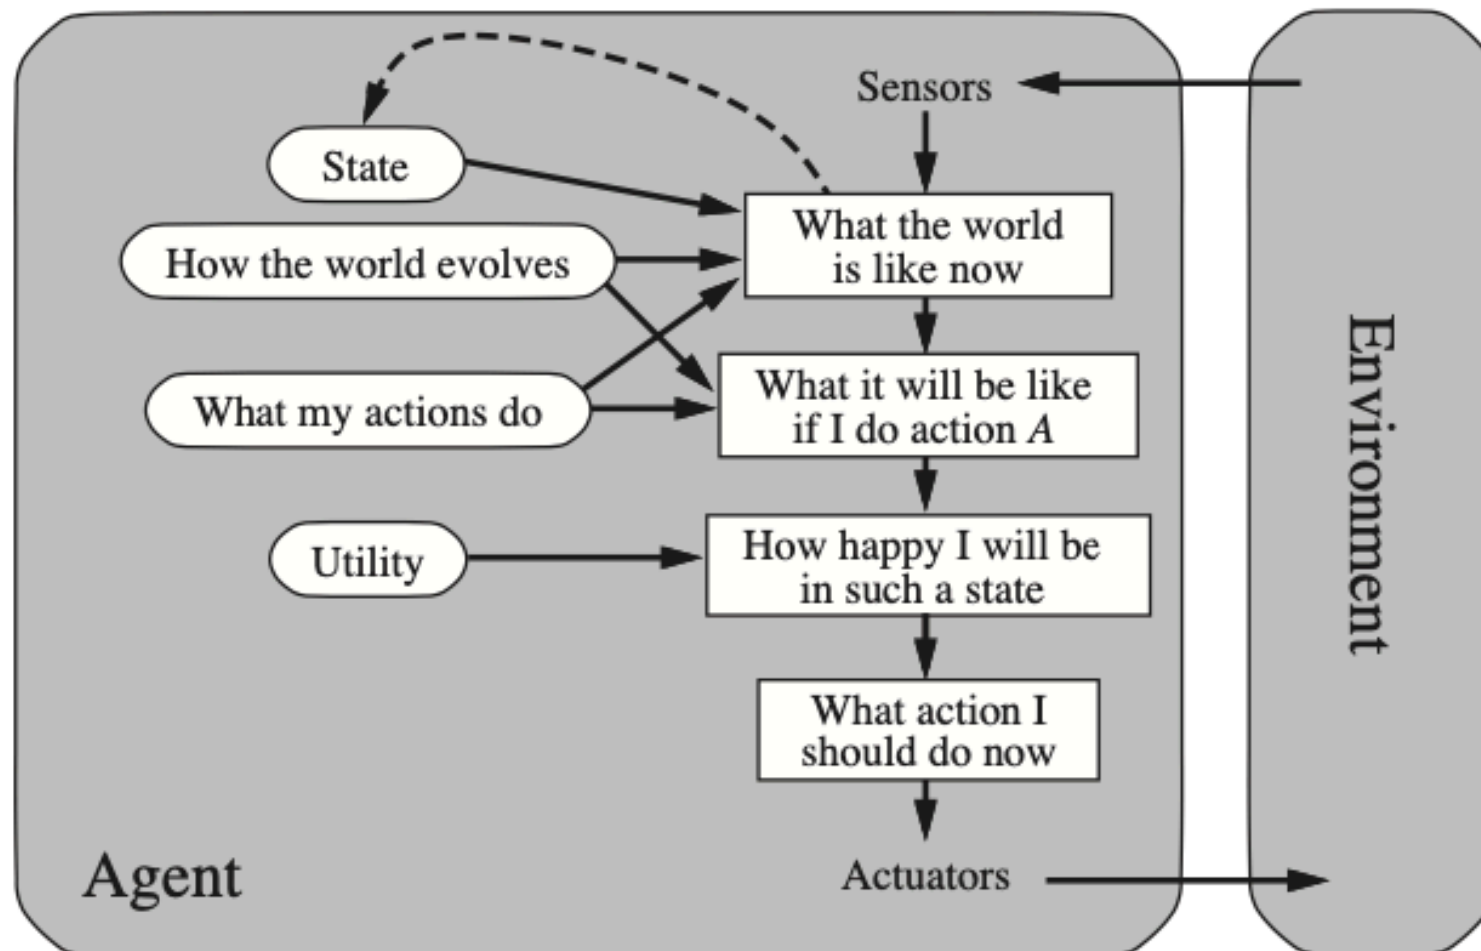


# Utility-based agent

- A rational utility-based agent chooses the action that maximises the **expected utility** of the action outcomes
  - that is, the utility the agent expects to derive, on average, given the probabilities and utilities of each outcome.
- The utility-based agent is not easy to implement
  - It has to model and keep track of its environment
  - Tasks involved a great deal of research on perception, representation, reasoning, and learning.
  - It can be implemented as a **Decision-making agent** that must handle the uncertainty inherent in stochastic or partially observable environments.

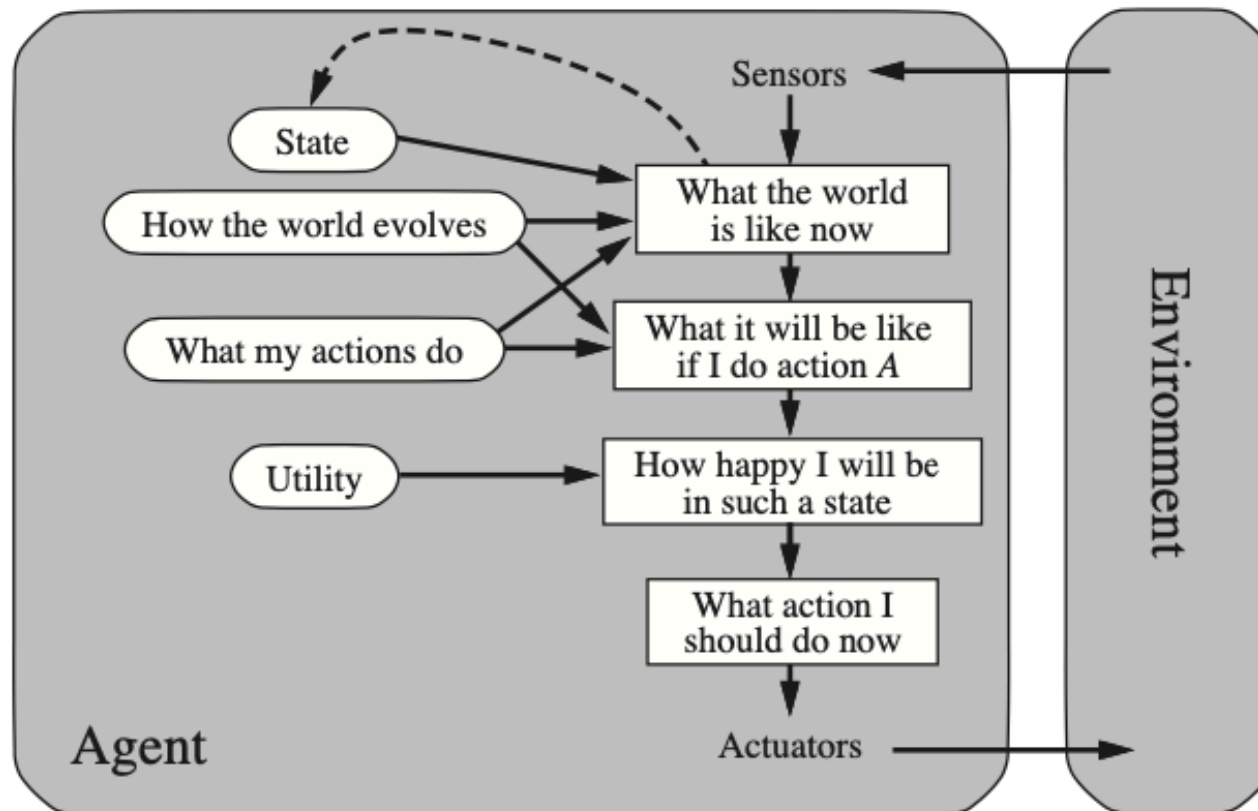


# Utility-based agent



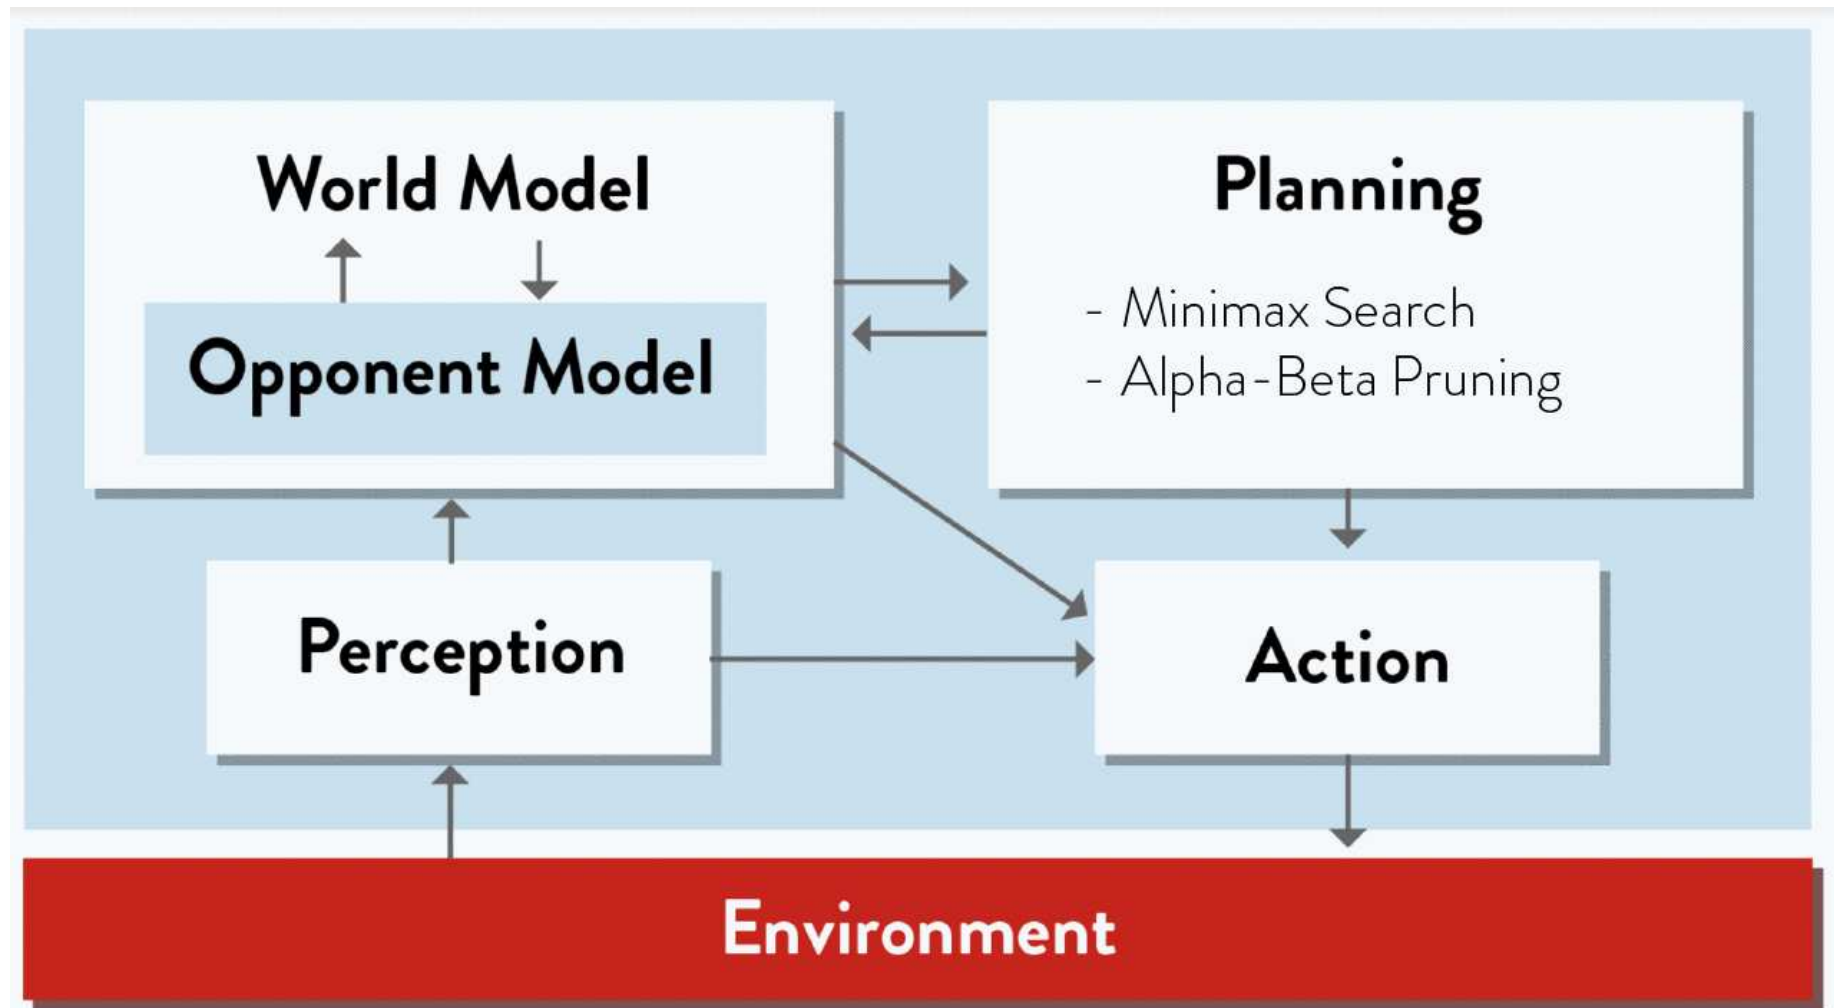
Utility-based agent — considers preference for certain world states over others

# Utility-based agent

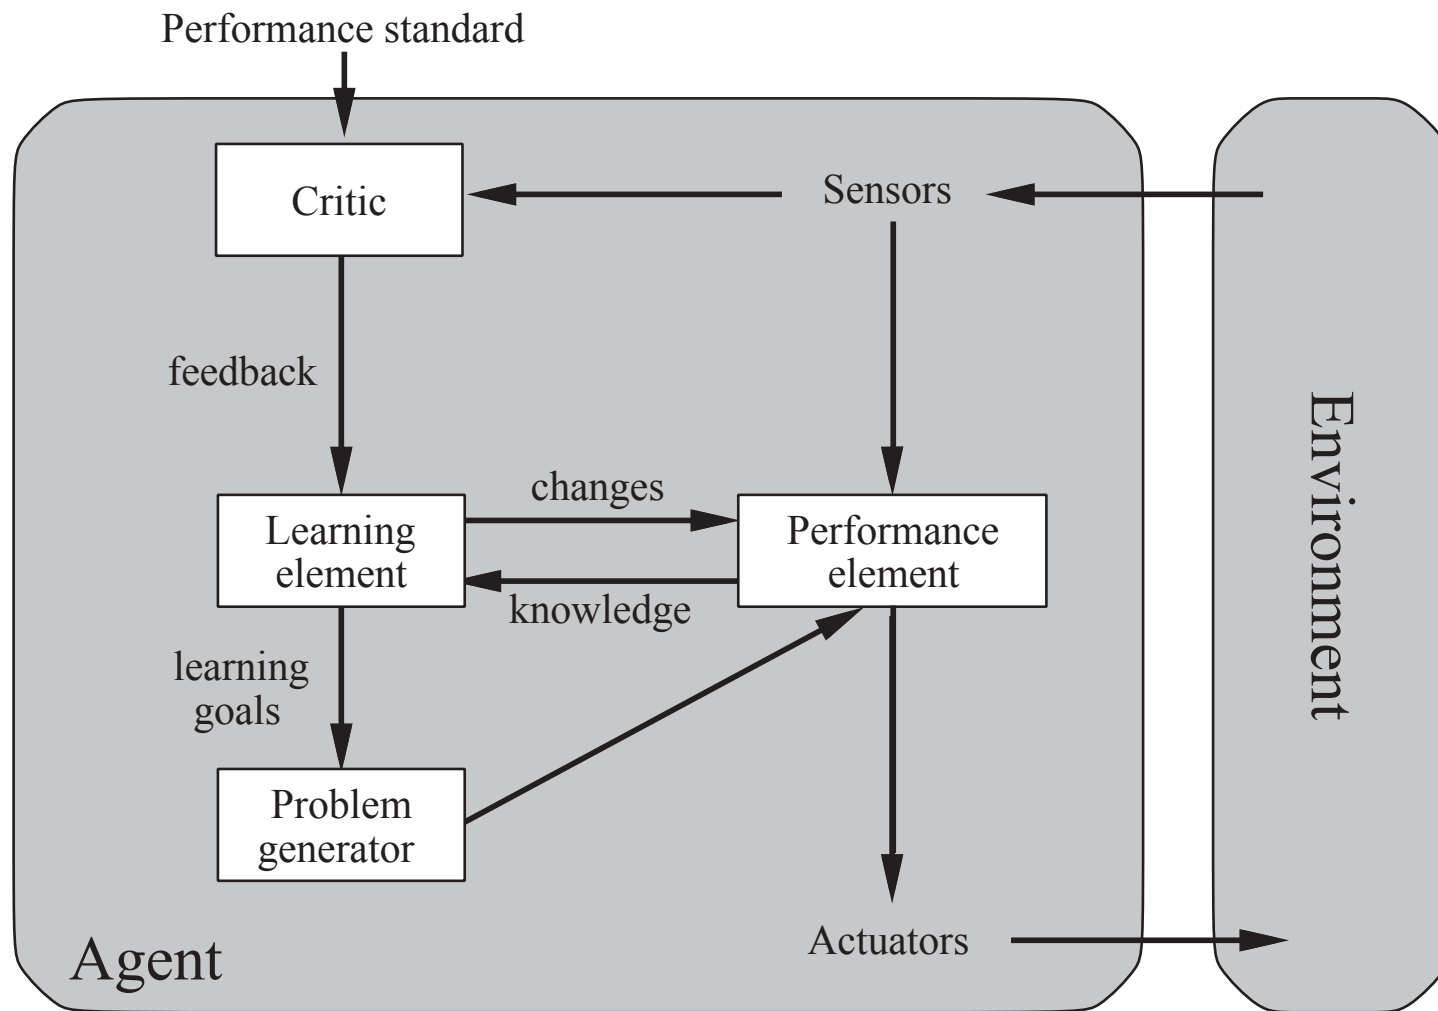


- A model-based, utility-based agent uses a model of the world, along with a utility function that measures its preferences among states of the world.
- It chooses the action that leads to the best expected utility, where expected utility is computed by averaging over all possible outcome states, weighted by the probability of the outcome.

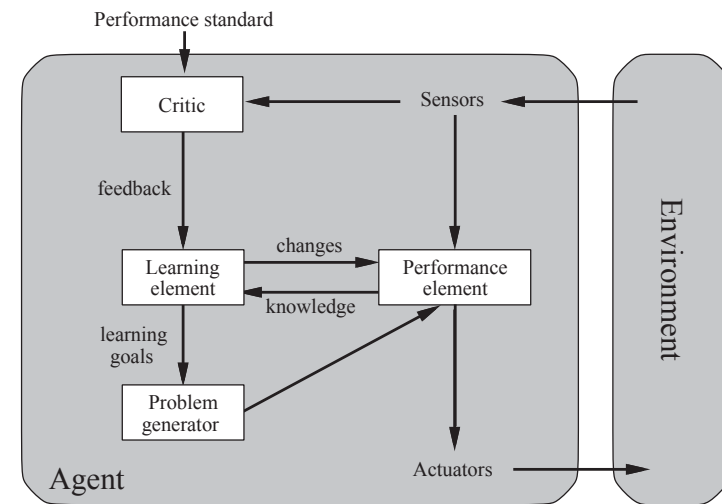
# Game Playing Agent



# Learning Agent

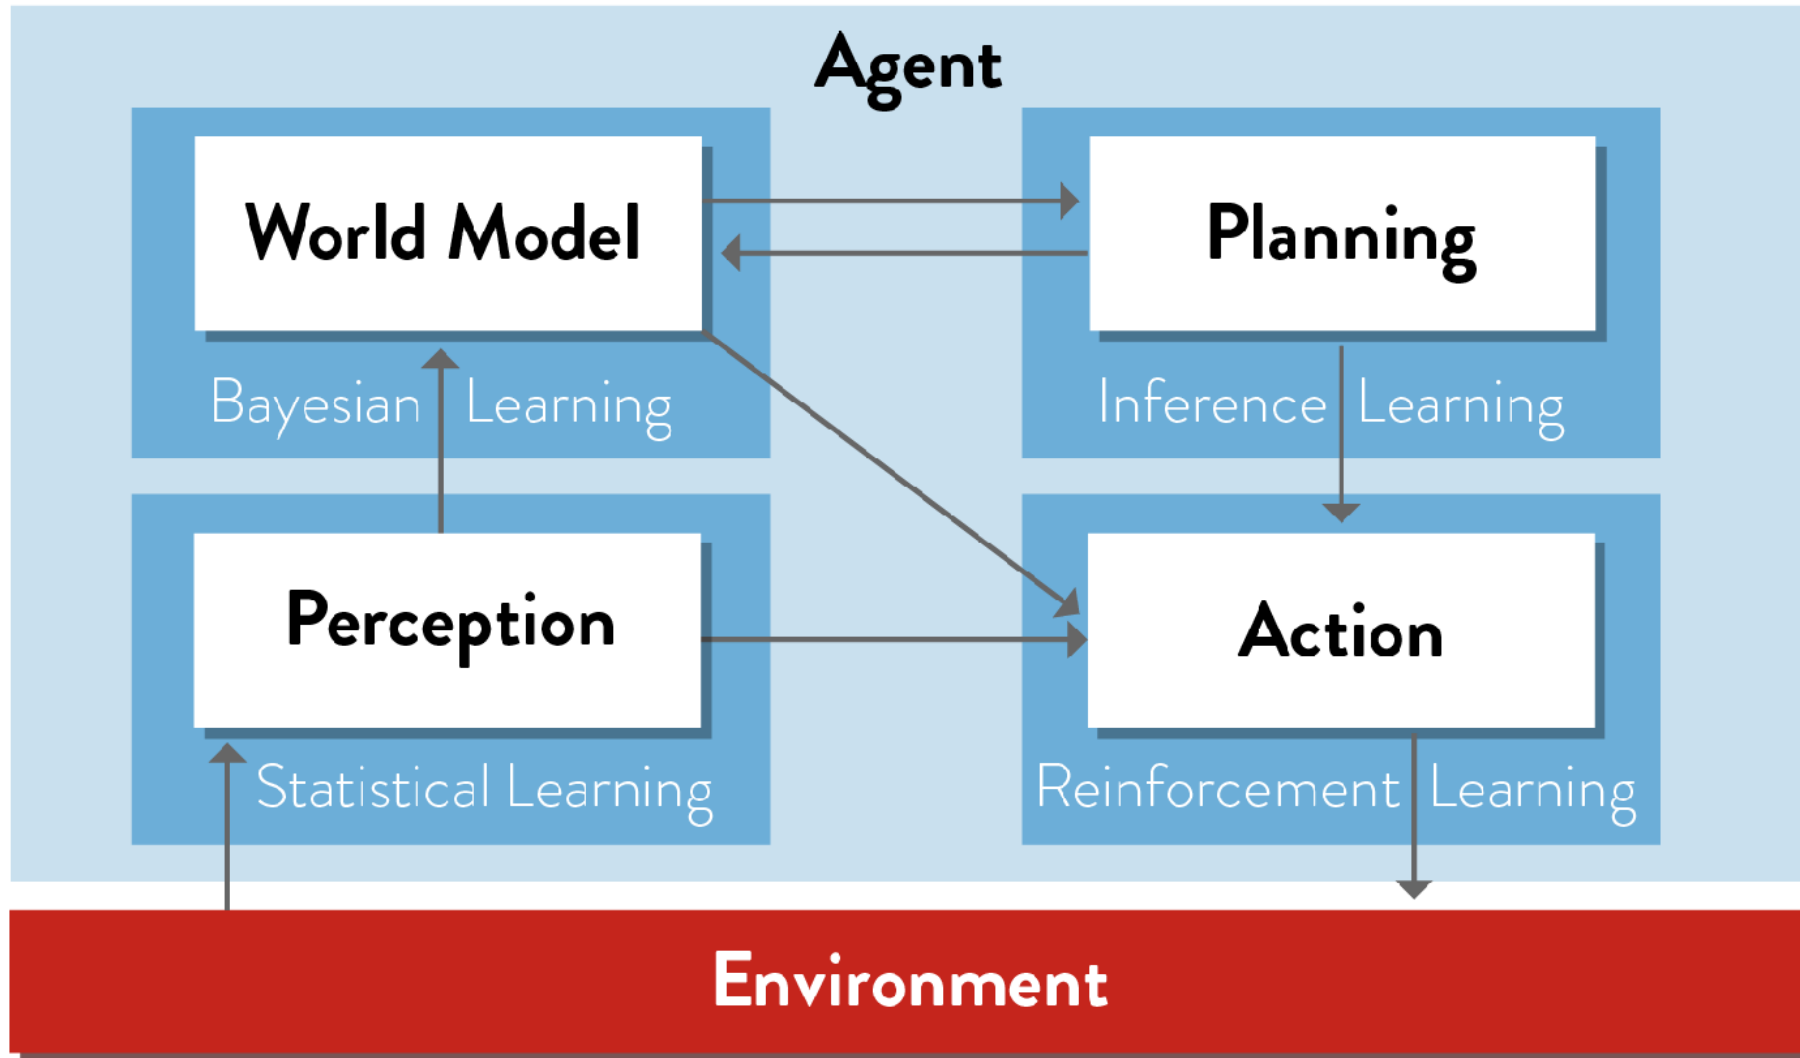


# Learning Agent



- The **performance element** takes in percepts and decides on actions.
- The **learning element** uses feedback from the **critic** on how the agent is doing and determines how the performance element should be modified to do better in the future.
- The **Problem generator** creates new tasks that provide new and informative experiences.

# Learning Agent



# Learning

- Learning is not a separate module, but rather a set of techniques for improving the existing modules
- Learning is necessary because:
  - may be difficult or even impossible for a human to design all aspects of the system by hand
  - the agent may need to adapt to new situations without being re-programmed by a human

# Summary


- Simple **reflex agents** respond directly to percepts,
- **Model-based reflex agents** maintain internal state to track aspects of the world that are not evident in the current percept.
- **Planning (Goal-base) agents** act to achieve their goals, and
- **Utility-based agents** try to maximise their own expected “happiness.”
- All agents can improve their performance through learning.



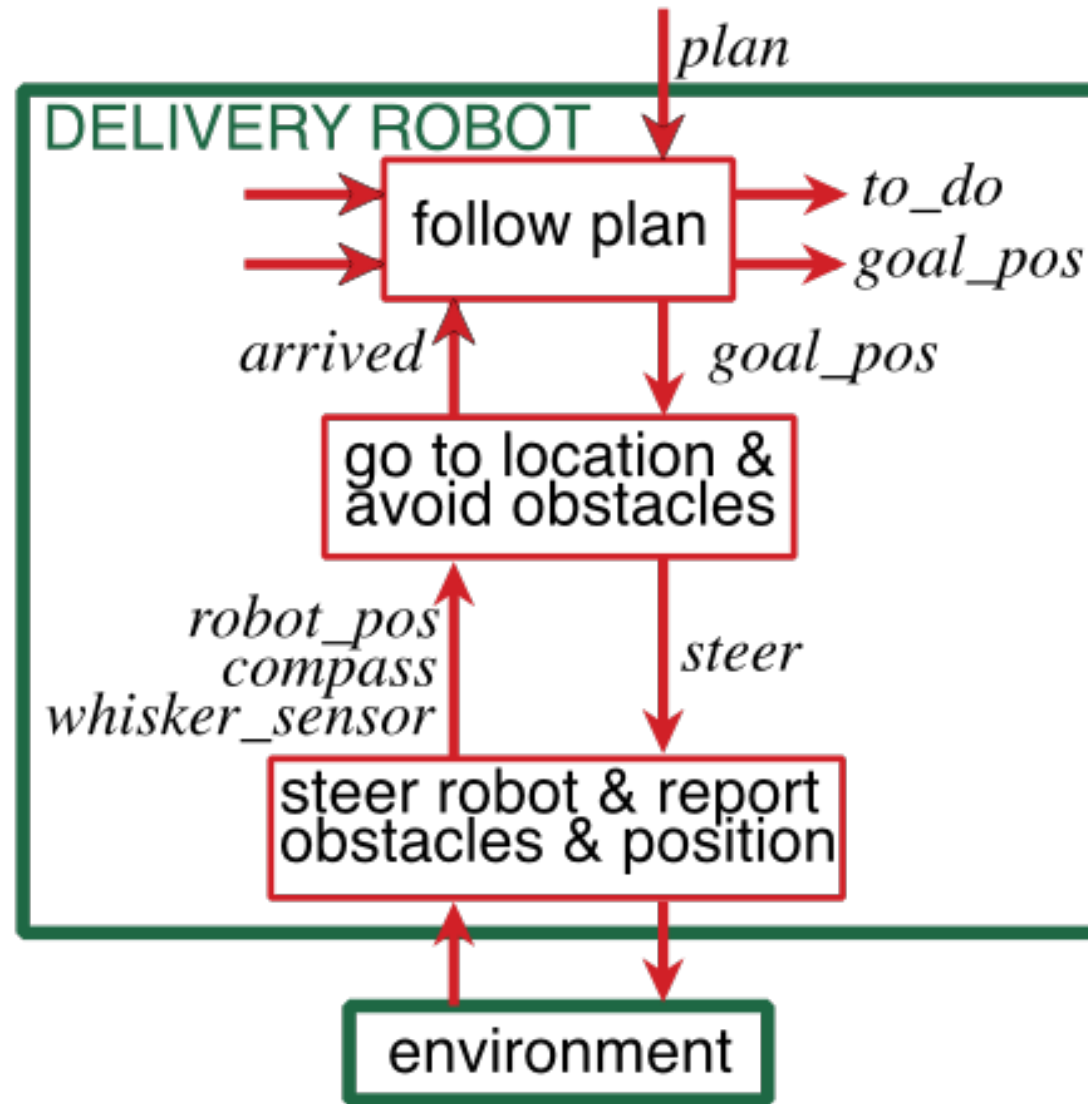
# Representation and Search

- The world model must be represented in a way that makes reasoning easy.
- Reasoning (problem solving and planning) in AI almost always involves some kind of search amongst possible solutions.

# Example – Delivery Robot



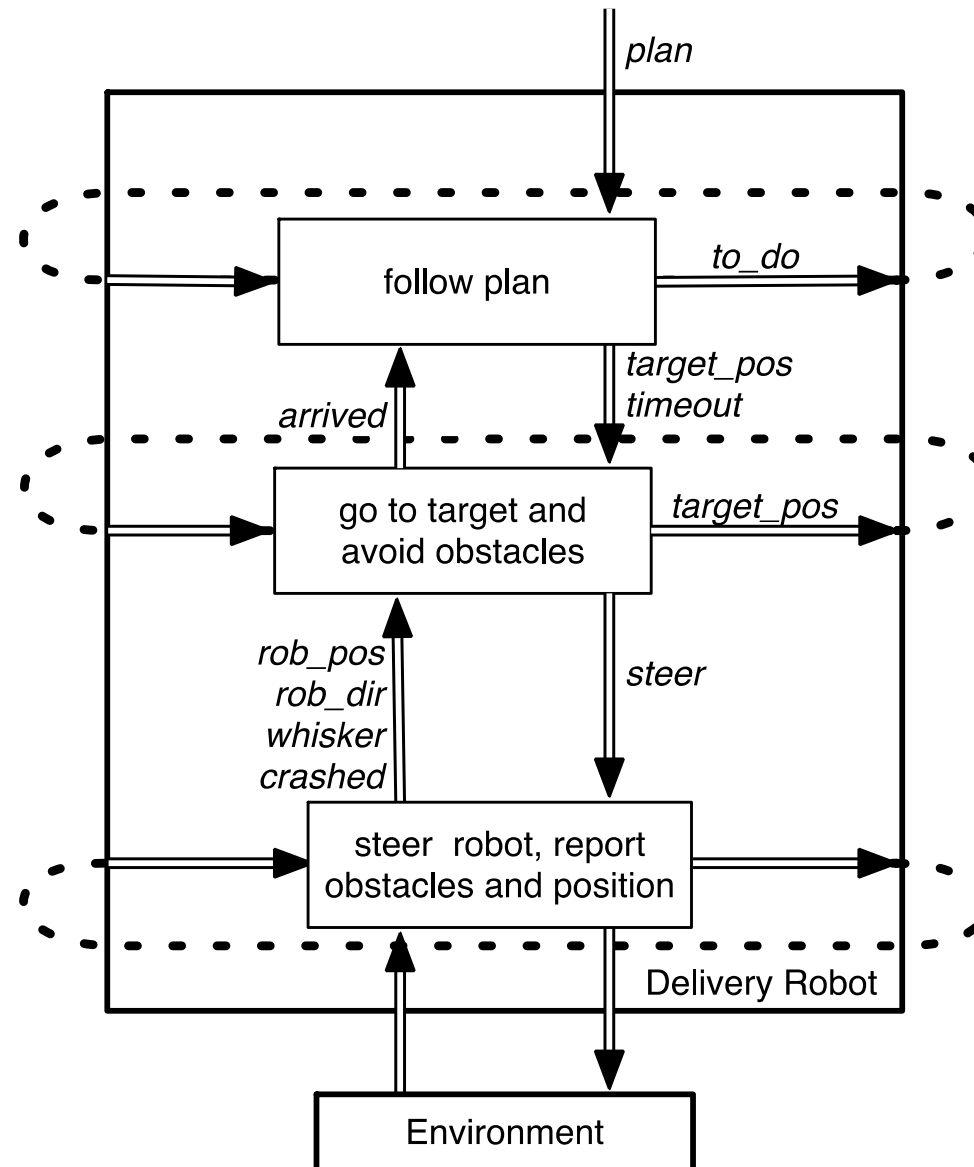
# Example – Delivery Robot



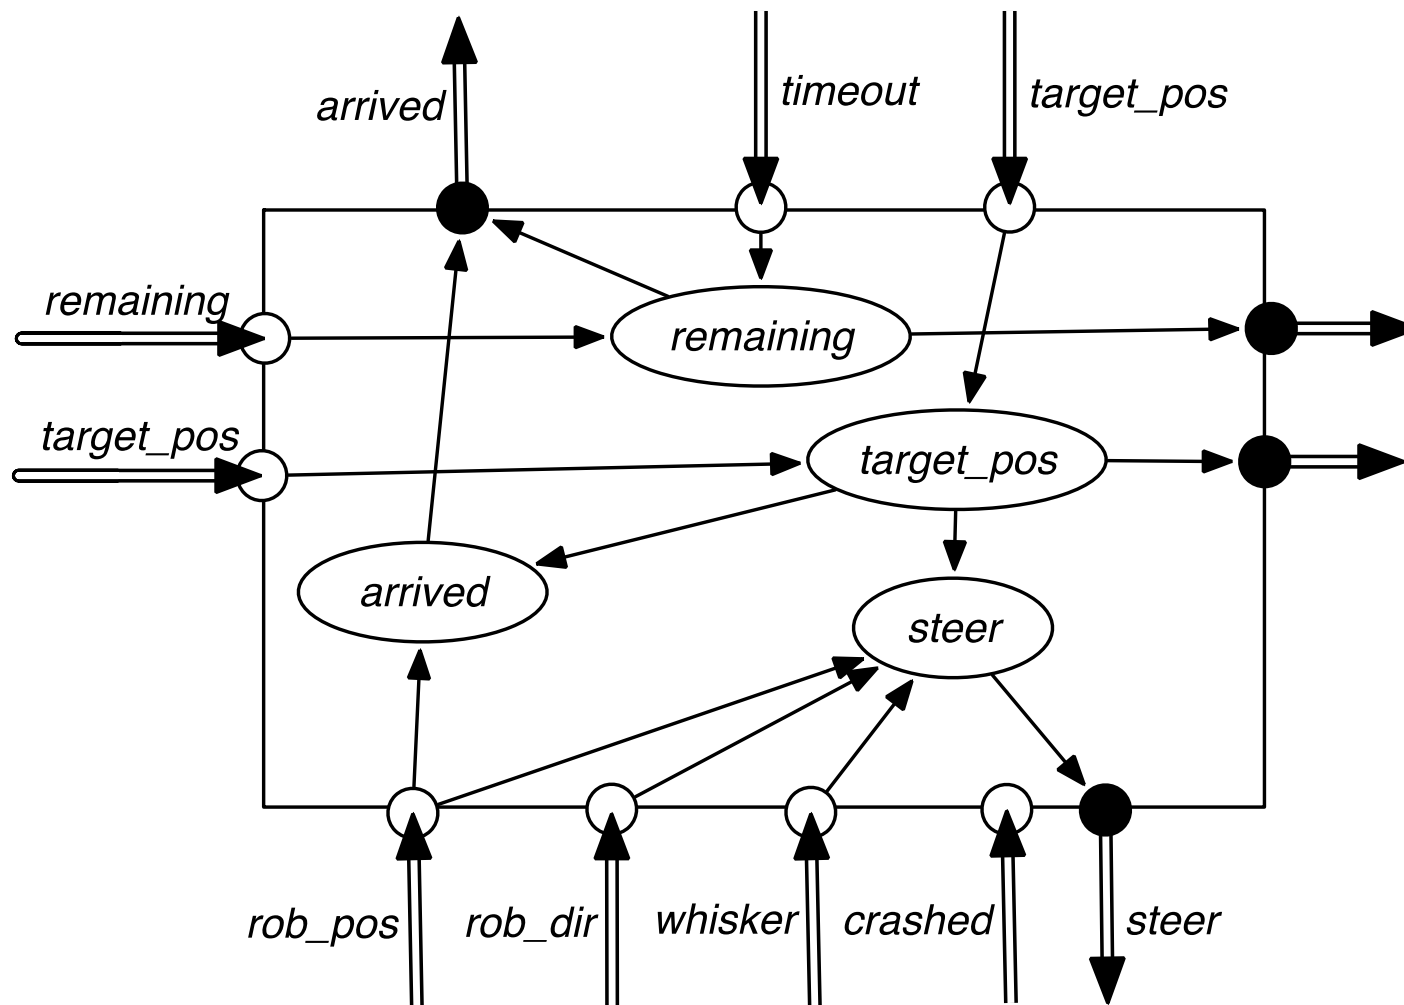
# Layered Architecture

- Hierarchy of controllers
- Controller gets percepts from and sends commands to the lower layer
  - Abstracts low level features into higher level (perception)
  - Translates high level commands into actuator instructions (action)
- The controllers have different representations, programs
- The controllers operate at different time scales
- A lower-level controller can override its commands

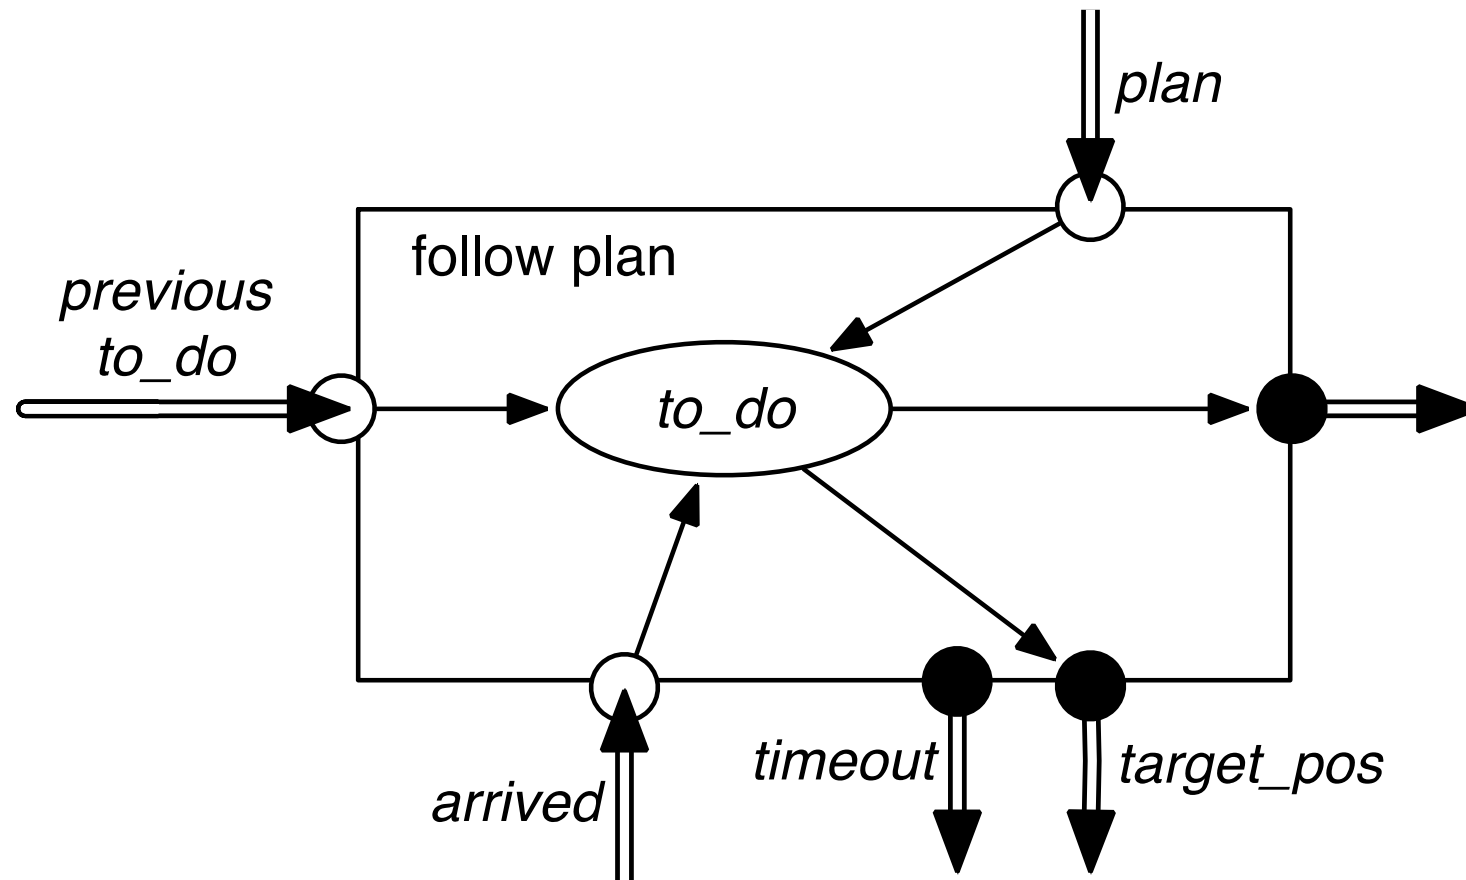
# Layered Architecture



# Delivery Robot – Middle Layer



# Delivery Robot – Top Layer



## Top Layer Code


```
given plan:
  to do := plan
  timeout := 200
  while not empty(to do)
    target pos := coordinates(first(to do))
    do(timeout; target pos)
  to do := rest(to do)
```

## Middle Layer Code

```
given timeout and target pos:
  remaining := timeout
  while not arrived() and remaining  $\neq$  0
    if whisker sensor = on
      then steer := left
    else
      if straight ahead(rob pos; robot dir; target pos)
        then steer := straight
      else
        if left of (rob pos; robot dir; target pos)
          then steer := left
        else steer := right
    do(steer)
  remaining := remaining - 1
  tell upper layer arrived()
```



# Delivery Robot – Simulation



# References

- Poole & Mackworth, *Artificial Intelligence: Foundations of Computational Agents*, Chapter 1 & 2
- Russell & Norvig, *Artificial Intelligence: a Modern Approach*, Chapter 2.